

Reaching the Actionable Radius: Physics-Guided Calibrated Localization of Methane Sources

Eric Liang

Department of Computer Science, Boston University
eric.jinchen.liang@gmail.com

Abstract—We present a simple recipe that brings calibrated machine-learning localization of high-rate methane sources to the 50 m scale used for facility association in EPA super-emitter notifications. On a methane-configured testbed whose closed-form physics makes the *exact* Bayesian posterior computable for every scenario, we first contribute an audit that measures, in meters, how much of the available information a calibrated localizer delivers: a conformal-wrapped deep localizer with a verified 90% coverage guarantee attains a 94 m median search radius where the exact posterior supports 6–7 m. Controlled elimination locates the cause: not the information, not the estimator head (a strengthened flow-based neural posterior estimator reaches only 63–71 m), and not the classical physics statistic alone (conformalized, it returns essentially the entire site) — but the absence of a classical computation, per-cell matched filtering, that pooled set encoders did not recover under the tested architectures and budgets. Supplying it directly — two closed-form matched-filter maps as input channels and a zero-initialized residual head, 0.2% extra parameters, ordinary label training — brings three architecture families into a 1.5 m band at 46–48 m median radius (a 50% median radius reduction; 96% of scenarios improved) at unchanged coverage and millisecond latency. Under realistic wind-measurement error, exact plug-in inversion retains its guarantee only by returning essentially the entire site, while the same recipe remains calibrated at 75–85 m, with uncertainty-aware input maps recovering a further ~ 8 m. Calibrated, physics-guided learning is both sharper and more robust than inversion built on the same physics.

Index Terms—methane monitoring, conformal prediction, simulation-based inference, sensor networks, physics-guided machine learning

I. INTRODUCTION

Methane delivers roughly $80\times$ the 20-year warming potency of CO_2 and decays within about a decade, so rapid repair of oil-and-gas *super-emitters* is among the highest-leverage near-term climate actions. Detection has advanced rapidly — UNEP’s satellite alert system issued over a thousand notifications in its first two years [1] — and attribution is a critical downstream bottleneck between detection and operator investigation. Regulation suggests a concrete scale: EPA super-emitter notifications associate alerts with facilities at a 50 m radius [2] (program implementation currently extended to 2027), with analogous source-level duties in the EU [3]. We adopt 50 m as an operationally motivated, facility-scale localization target — a region that triggers a focused investigation, not yet a repair.

Continuous monitoring with sparse fixed point-sensor networks is where such localization can happen: a handful of cheap masts around a facility, always on, catching the plume as it sweeps across them under naturally meandering wind. Machine learning supplies millisecond inference, and split conformal prediction [4] supplies a guarantee: a region provably containing the source 90% of the time. **The gap:** coverage guarantees frequency, not usefulness. Two questions had no quantitative answer: (i) how close is a calibrated region to the smallest region the sensor data support under the modeling assumptions, and (ii) what mechanism limits it, and what closes it? Standard diagnostics for amortized inference test whether a posterior is calibrated [5]; they cannot measure its distance from the exact posterior at matched calibration.

Contributions.

- **C1 — An exact-posterior audit.** A methane testbed (CH₄-T) whose closed-form physics makes the true Bayesian posterior of every scenario computable, so conformal efficiency can be measured in meters against a model-based information reference (6.2 m at matched empirical coverage) and the 50 m facility scale.
- **C2 — A diagnosis by elimination.** Controlled comparisons — a labels-only baseline, a strengthened flow-based neural posterior estimator, and the conformalized classical statistic without any network — isolate the dominant bottleneck in our experiments: a classical computation (per-cell matched filtering) that pooled encoders did not recover, and that neither a modern estimator head nor the raw physics statistic supplies alone.
- **C3 — A minimal recipe that reaches the scale.** Two closed-form matched-filter maps plus a zero-initialized residual head (0.2% extra parameters, ordinary label training) bring three architecture families to 46–48 m median radius at unchanged coverage and latency, and remain calibrated under wind-measurement error that reduces exact plug-in inversion to site-sized regions.

The components are deliberately classical — matched filters are 1940s detection theory; conformal prediction and label training are standard — and the composition is minimal: the ablation record (Appendix) shows both map channels contribute and neither the network nor the statistic suffices alone. The exact posterior’s role throughout is measurement,

Figure 1 placeholder: wind-swept plume, masts, and true source on the 500 m site, with median 90% regions drawn to scale — exact-posterior reference (~ 7 m), our recipe (46–48 m), labels-only baseline (94 m), and the 50 m facility-association radius.

Fig. 1. The opportunity, to scale (placeholder).

not supervision.

II. RELATED WORK

Conformal prediction provides distribution-free coverage [4]; efficiency is the open half of the problem, and we contribute an audit that scores it against the exact Bayesian reference at matched, verified coverage. **Simulation-based inference** amortizes posteriors with neural density estimators [6], [7] at simulation budgets matching ours [8]; its diagnostics test calibration [5], and our audit adds the efficiency question while placing a representative flow-based estimator inside the comparison. **Source-term estimation** inverts dispersion models via optimization or MCMC at minutes-to-hours per event [9]; deep localizers give millisecond inference but lacked instruments certifying how much information they extract. **Physics-guided ML** injects domain structure through features or architectures; our input maps are the forward model’s sufficient statistics under CH4-T assumptions — classical matched-filter detection theory supplied as spatial channels aligned with the output grid. **Controlled-release evaluation** (METEC) scores continuous-monitoring systems on detection, quantification, and equipment-scale localization [10]; CH4-T mirrors that task structure at super-emitter rates with exact ground truth, complementing — not replacing — field validation.

III. METHODOLOGY

A. Testbed and exact-posterior audit (C1)

CH4-T places a source of unknown location and rate ($10\text{--}500\text{ kg h}^{-1}$, log-uniform, bracketing the 100 kg h^{-1} super-emitter threshold) on a 500×500 m site sensed by 4–12 masts (ppm readings, 0.2–3 ppm noise, dropout). The plume is a steady Gaussian under known Pasquill–Gifford stability with Briggs open-country dispersion (coefficients verified against published tables; mass flux conserved to 4×10^{-6}), driven quasi-statically by a known fluctuating wind record ($1\text{--}8\text{ m s}^{-1}$, AR(1) meander; steps read as 1-min averages, a 30-min window). Because the response is closed-form and linear in rate, the exact posterior over a 64×64 grid is computable per scenario (adaptive rate quadrature, float64, normalization error $\sim 10^{-12}$). A blocking gate requires conformal prediction applied to the exact posterior itself to self-calibrate — a gate

that caught two implementation bugs, including a tie-handling error at degenerate scores whose fix we report for reuse. All models share identical data (10k/1k/2k/2k splits, the modal SBI budget [8]), D4 symmetry augmentation, and an identical randomized-HPD split-conformal wrapper at verified 90% coverage. Regions are reported as equal-area circle radii; the exact-posterior reference is 7.0 m at nominal 90% (empirical 0.95) and 6.2 m at matched empirical 0.89.

B. Diagnosis by elimination (C2)

A DeepSets localizer trained on smoothed one-hot labels is calibrated (coverage 0.89) at a 94 m median radius — the information for 6–7 m is in the readings, waiting to be extracted. Three controlled probes isolate why. *Is it the estimator head?* A representative flow-based neural posterior estimator on the same encoder — strengthened after adversarial review of our own implementation ($2.5\times$ budget, widened head, full-band conditioning) — reaches 62.8–71.2 m, and we verify its spline head can express 4.4 m regions, so its remaining blur is not head-structural. *Is the classical physics statistic enough by itself?* Conformalizing the normalized matched-filter map with no network returns essentially the entire site: the statistic is overconfident-wrong, and calibration must be learned. *Is the computation the missing piece?* Localization requires correlating each candidate cell’s space–time signature against all readings — per-cell matched filtering — which pooled set representations did not recover under the tested architectures and budgets; supplying exactly this computation as input (§III-C) collapses the gap on every encoder. Together: the network cannot compute what the statistic knows; the statistic cannot calibrate what the network learns; the combination is the method.

C. The recipe (C3)

For any base localizer (Fig. 2): **(1)** compute per cell c the sufficient statistics $a_c = \sum g_c y$, $b_c = \sum g_c^2$ of the linear-Gaussian model (g_c the unit-rate response) and feed two 64×64 channels: the matched-filter statistic $z_c = a_c / (\sigma \sqrt{b_c})$ and $\log b_c$ — milliseconds, closed form, computed from the network’s own inputs (no new information, only a computation); **(2)** add a 3-layer zero-initialized convolutional head (~ 10 k parameters, 0.2% of the base) summed with the localizer’s logits, so at step 0 the model is the unmodified baseline and every comparison is controlled by construction; **(3)** train with ordinary smoothed-label cross-entropy; **(4)** wrap with unchanged split conformal. Physics computes the evidence; the network weighs it honestly; conformal certifies the result.

Why not invert the physics directly? On this testbed one can — the exact posterior is our measuring instrument, tractable by construction. In deployment, realistic dispersion has no closed form and simulator inversion costs minutes-to-hours per event [9] where monitoring needs seconds — and, as §IV-C shows, plug-in inversion is brittle under wind-model misspecification in a way the learned recipe is not.

Figure 2 placeholder (full width): method pipeline. Observed readings + wind record $\rightarrow$ matched-filter maps $z, \log b$ (closed form, ms) $\rightarrow$ any localizer + zero-initialized conv head, joined at the logits $\rightarrow$ split-conformal 90% region (median 94 m $\rightarrow$ 46–48 m against the 50 m facility scale).

Fig. 2. Physics-guided calibrated localization (placeholder).

TABLE I

MEDIAN 90% CONFORMAL REGION RADIUS (M); RANGE OVER 3 SEEDS; COVERAGE 0.89–0.92 IN EVERY CELL. REFERENCES: EXACT-POSTERIOR 6.2–7.0 M, FACILITY-ASSOCIATION SCALE 50 M.

Encoder	Labels only	+ Physics maps
DeepSets	94–98	47.1–47.7
GNN	55–58	46.2–47.5
Set Transformer	69–75	46.2–47.3

TABLE II

A REPRESENTATIVE FLOW-BASED NPE HEAD ON THE SAME ENCODER, DATA, AND AUDIT; 3 SEEDS PER ROW; ALL LABEL-TRAINED.

Model (DeepSets encoder)	Radius (m)	Coverage
Flow-NPE (strengthened)	62.8–71.2	0.90–0.91
Flow-NPE + physics maps	57.6–61.2	0.89–0.91
Ours (grid head + maps)	47.1–47.7	0.89

IV. RESULTS

A. One recipe, three architectures

The maps improve every architecture, and the improvement is uniform to a striking degree: without them, encoder choice spans 40 m (94 vs. 55); with them, all nine seed-runs land inside a 1.5 m band (Table I) — the gains are available to whatever architecture a practitioner already has, and encoder design ceases to be the binding constraint. Against the labels-only baseline the median radius reduction is 50.3% (bootstrap 95% CI [48.9, 52.0]; 96.4% of scenarios improved). Point accuracy is essentially unchanged: the maps improve the *shape of uncertainty*, which region size measures. Operationally: 52% of individual test cases fall below 50 m (median 47 m); the 75th/90th/95th percentiles are 86/131/146 m; performance is mildly worse at the highest rates (56.9 m median at 250–500 kg h⁻¹, vs. 42–43 m below 100 kg h⁻¹) — the median case reaches the facility scale, and the upper tail defines the remaining challenge.

B. Comparison with a modern flow-based estimator

The flow head is a genuine lever — it improves markedly on the categorical labels-only baseline with no physics at all — and the maps transfer onto it (Table II). The convolutional head reading the maps *spatially* extracts more than the flow’s parametric conditioning at identical information. Our approach is compared against, and complements, a modern estimator rather than being superseded by it.

TABLE III

WIND MEASURED WITH REALISTIC ERROR (10° DIRECTION, 10% SPEED); TRAINING, CALIBRATION AND TEST ALL UNDER THE SAME ERROR DISTRIBUTION; 3 SEEDS FOR LEARNED ROWS.

Pipeline	Coverage	Radius
Plug-in inversion (measured wind)	1.00	282 m (full site)
Plug-in, inflated-noise robust variant	1.00	282 m
Ours, deterministic maps	0.91–0.92	81.7–84.9 m
Ours, ensemble maps	0.89–0.90	74.7–75.4 m

C. Robust where plug-in inversion is not

Real anemometers err. Under a 10°/10% observation error, plug-in inversion at the measured wind assigns numerically negligible (below float64) posterior mass to the true source in 30% of scenarios; split conformal preserves its guarantee exactly as theory requires — by returning essentially the entire site (Table III). A global noise-inflation robust variant, tuned on calibration data, does not help: per-step wind error is not absorbable by a scalar noise scale. The recipe, trained through the same error, remains calibrated at useful sizes, and *ensemble maps* — the mean and spread of the matched-filter maps over eight draws from the wind-error model — recover a further ~ 8 m. Under the modeled wind-measurement error, learned physics-guided localization is more robust than plug-in inversion built on the same physics; a wind-marginalized exact reference is future work.

V. LIMITATIONS

CH4-T is an instrument, not a deployment: single steady sources (no zero-leak or multi-source scenarios), flat terrain without equipment wakes, known stability class, quasi-static stepping, in-distribution marginal coverage, and a 64×64 grid whose ~ 7.8 m cells sit near the oracle reference scale (a grid-refinement study is pending). The 6–7 m reference is a property of CH4-T’s generative assumptions, not an atmospheric limit. Half of individual cases still exceed 50 m, and the upper tail (p90 = 131 m) is unresolved. The architectural claim is scoped to tested encoders and budgets. Field claims await controlled-release validation — comparisons here are exact regardless, because all models share one verified simulator.

VI. CONCLUSIONS

On a testbed with exact Bayesian ground truth, a calibrated deep localizer delivered a tenth of the sharpness its sensor data support; controlled elimination traced the gap to a single

classical computation, and supplying that computation — two matched-filter maps and a 0.2%-parameter zero-initialized head, ordinary labels, nothing else — brought three architectures into a 1.5 m band at the 50 m facility-association scale, halving the median search radius at unchanged cost, ahead of a strengthened flow-based estimator, and calibrated under wind error that reduces exact inversion to site-sized regions. The constructive message: simulators make exact posteriors available as *instruments* for auditing what learned systems extract, and classical signal processing still knows computations that networks need not rediscover. We hope the evidence standard — exact-posterior audits, blocking self-calibration gates, one-change-at-a-time controls — travels furthest.

APPENDIX

APPENDIX: ABLATION AND CONTROL RECORD

DeepSets encoder, identical data and conformal harness, median 90% region radius (m); ^{3s} marks 3-seed ranges.

TABLE IV
COMPONENT ABLATIONS (CLEAN-WIND TRACK), ALL LABEL-TRAINED.

Configuration	Radius (m)	Coverage
Labels only (baseline)	94–98 ^{3s}	0.89
Maps + head (the recipe)	47.1–47.7^{3s}	0.89
z channel only	65.0	0.88
$\log b$ channel only	98.5	0.90
Classical statistic only (no network)	282 (full site)	1.00
Exact-posterior reference	6.2–7.0	0.89–0.95

Reading. Both channels contribute and the matched filter dominates; the sensitivity channel alone ($\log b$, no correlation with the readings) adds little over the baseline, confirming the correlation computation is the operative ingredient. The classical statistic without the network is overconfident-wrong and pays with full-site regions — the learned calibration is load-bearing. Rate-stratified performance of the recipe (median radius / fraction below 50 m): 10–30 kg h^{−1}: 42.0 m / 55%; 30–100: 43.4 m / 58%; 100–250: 50.1 m / 50%; 250–500: 56.9 m / 41%.

REFERENCES

- [1] UN Environment Programme (IMEO), “An eye on methane: Methane Alert and Response System annual report,” 2024.
- [2] U.S. EPA, “Standards of performance for GHG emissions: NSPS OOOOb/EG OOOOc and the Super-Emitter Program,” 89 FR 16820, 2024.
- [3] European Union, “Regulation (EU) 2024/1787 on methane emissions reductions in the energy sector,” OJ L, 2024.
- [4] A. N. Angelopoulos and S. Bates, “Conformal prediction: A gentle introduction,” *Found. Trends Mach. Learn.*, 16(4), 2023.
- [5] S. Talts *et al.*, “Validating Bayesian inference algorithms with simulation-based calibration,” arXiv:1804.06788, 2018.
- [6] G. Papamakarios and I. Murray, “Fast ϵ -free inference of simulation models with Bayesian conditional density estimation,” *NeurIPS*, 2016.
- [7] C. Durkan, A. Bekasov, I. Murray, and G. Papamakarios, “Neural spline flows,” *NeurIPS*, 2019.
- [8] J.-M. Lueckmann *et al.*, “Benchmarking simulation-based inference,” *AISTATS*, 2021.
- [9] T. Newman *et al.*, “Probabilistic inversion modeling of gas emissions: MCMC estimation of Gaussian plume parameters,” arXiv:2408.01298, 2024.

- [10] C. Ilonze *et al.*, “Assessing the performance of emerging and existing continuous monitoring solutions under a single-blind controlled testing protocol,” *Elementa*, 2025.